

Introduction to Programming with R

Dries Debeer & Benjamin Becker

31. September and 01. October 2021

Zurich R Courses

Table of Content

Introduction	if
Motivation	ifelse
Goal	Loops & Iteration
Agenda	for
RStudio setup	while & repeat
Basic elements & data types	next & break
What are objects?	Loops & Iteration - good practice
Atomic Vectors	Functionals
Vector Structures	Functions I
Subsetting - Selection	Function Calls
Element Replacement	Function Names
Flow & conditional programming	Arguments
	Output

Introduction

Introduction

Who are we?

Dries Debeer

Senior Researcher at itec
(imec Research Group at KU
Leuven)

scDIFtest, permimp, eatATA,
mstDIF

dries.debeer@kuleuven.be

Benjamin Becker

Researcher at IQB (Statistics
Department)

eatGADS, eatDB, eatATA,
pisaRT

b.becker@iqb.hu-berlin.de

Who are you?

1. Occupation, employer?
2. Previous knowledge and experience
 - with R?
 - with other statistical software?
 - with other programming languages?
3. Specific interest/motivation for this workshop?

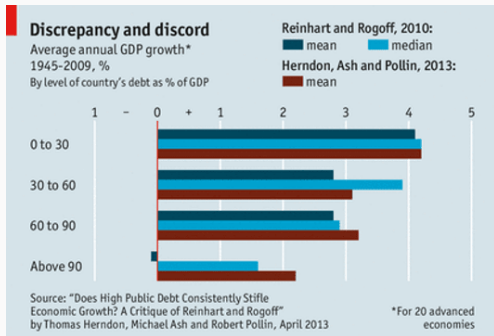
Motivation

1. Increase efficiency!
 - Save time and nerves
 - Avoid errors and bugs
 - High transfer effect to all projects (with data analyses)
2. Successful collaborations (including with your future self!)
3. Code as deliverable (i.e., part of research paper)

Two of your worst collaborators

- Past Self
 - the biggest mess in existence
 - Did not document anything
 - Uses a completely different style of writing code
 - does not reply to e-mails
- Future Self
 - Has the memory of a goldfish
 - Will have zero understanding for your current brilliance

Motivation



The Sunday Telegraph Sunday 17 May 2020

Coronavirus

Telling behind lockdown was an unreliable buggy mess, claim experts

Data that predicted 500,000 could die in UK unless extreme measures were taken are impossible to replicate, say scientific teams

Science

By Hannah Ireland
and Tilly Zaidenphart

THE Covid-19 modelling that sent Britain into lockdown, sending the economy and leaving millions out of work has been criticised by experts.

Prof Neil Ferguson's Imperial College computer coding was derided as "totally unreliable" by leading figures, who claimed it was "unsure if you wouldn't shake your fist at it".

The model, created with funding from the Government to U-turn and introduce a nationwide lockdown, is a "buggy mess, which looks more like a bowl of spaghetti than a finely tuned piece of programming", said David Foray, the co-founder of data technology company WANdisa.

"If we managed to build a model, we would have anyone for developing code the this and any business that relied on it to produce software for sale would then go bust."

The comments are likely to relegate a row over whether the UK was right to go into lockdown, with conflicting models suggesting people may have been ready acquired substantial lost immunity and Covid-19 may have hit Britain earlier than first thought.

Scientists have also been split on the fatality rate of Covid-19, which has resulted in vastly different models.

Up until now, significant weight has been attached to Ferguson's model, which placed the fatality rate higher than others and predicted 500,000 in the UK could die without a lockdown.

It was said to have prompted a dramatic change in government policy, closing businesses, schools and restaurants to be shut immediately in March. The Bank of England has predicted that the economy could take a year to return to normal, after its worst recession in more than three centuries.

The Imperial model works by using code to simulate transport links, population size, social networks and knowledge gaps. It provides a prediction for how many people would be infected over time. However, questions have emerged over whether the models accurately, after researchers realised its code, which is in the public domain was "horrendous at times" developed over more than 10 years.

In its initial form the code was unreliable, developers claimed, with some parts looking "like they were written

'In our commercial reality, we would fire anyone for developing code like this'

'Any business that relied on it to produce software for sale would likely go bust'

'It looks more like a bowl of angel hair pasta than a finely tuned piece of programming'

'The early 2000s were yet another confirmation that their modelling approach was flawed to the core'

code translated from Fortran", an old coding language, according to John Cunniff, a US developer, who said: "More than the code before it was published. Yet, the problems appear to go much deeper than messy coding. Many have claimed that it is almost impossible to reproduce the same results from the same data, using the same code."

Scientists from the University of Edinburgh said they got different results when they used different data

sets, and even in some cases using the same machines. "There appears to be a bug in either the creation or re-use of the network file. It's not a simple bug, it's a complex one, only varying in that the second should use the first, while the first should use the second, and the results are quite different," the Edinburgh researchers wrote on the Twitter social network. After a discussion with a GitHub developer, it was pointed out that it is said to be one of a number of

bugs discovered within the system. GitHub developers said that the model was "stochastic" (random), and "multi-parallel" with different results depending on the order of the code. It has prompted questions from scientists, who say "models must be capable of passing the basic scientific test of providing the same results given the same initial set of parameters" - otherwise, there is simply no way of knowing whether they will be reliable."

experts' predictions under different assumptions, and with different interventions, in increasing powers."

Like the Imperial code, a rival model by Prof Steven Caletka of the University of Oxford works on a so-called "SIR approach" in which the population is divided into those that are susceptible, infected and recovered. However, while Prof Caletka assumed that 0.1 per cent of infected people would die, Prof Ferguson worked on 0.1 per cent. That led to a dramatic reversal in government policy from attempting to limit "herd immunity" to a full-on lockdown.

Concerns over Prof Ferguson's model have now resurfaced, with Dr Kenneth Leung, the VP of architecture at WANdisa, saying his track record did not inspire confidence. In the early 2000s, Prof Ferguson's models incorrectly predicted up to 100,000 and 100,000 deaths, predictions from 100,000 and 100,000 from 100,000.

"The fact from the early 2000s are just yet another confirmation that their modelling approach was flawed to the core," says Dr Leung. "We don't know for sure if the same model was used, but we clearly see their modelling wasn't rigorous and that surely hasn't improved now."

A spokesman for Ferguson's team said: "The Government has never relied on a single disease model to inform decision making. As has been repeatedly stated, decision makers around lockdown was based on a consensus view of the scientific evidence including several modelling studies by different academic groups."

Multiple groups using different models concluded that the pandemic would overwhelm the NHS and cause unacceptable high mortality in the absence of extreme social distancing measures. When the Imperial team used we several models of differing levels of complexity, all of which produced consistent results. We are working with leading academic groups and technology companies to develop, test and refine the same model. The fatality rate referred to, however, we reject the further extension of the model to biologically-mediated communications.

"Epidemiology is not a branch of computer science and the consensus around lockdown was based on the scientific consensus that Covid-19 is highly transmissible with an infectious fatality rate exceeding 0.1 per cent in the UK."



Motivation

Concept of Technical Debt

- We write (messy) code for data cleaning/analyses
- We decide on data sets/models/graphs/tables/...
- We try to publish it, get a major revision
- We need to rerun some analyses
- Modifying/extending our code is more difficult than it should be

Trade-off

- Being fast vs. writing (or refactoring) perfect code

But also

- Write better R code

An introduction to R as a Programming language

- Better practical R skills
- Better theoretical understanding of R (and programming)
- Different framing: R as a programming language

Agenda

Day 1

- RStudio setup
- Basic elements & data types of the R language
- Flow & conditional programming
- Loops & iteration
- Writing & using functions (part I)

Day 2

- Writing & using functions (part II)
- Debugging
- Good programming practices

RStudio setup

RStudio setup

1. Copy the course content from the usb-stick to a directory on your machine
2. Open RStudio
3. Choose `File < New Project ...`
4. Choose `Existing Directory`
5. Browse to the directory on your machine where you copied the course content and select the “Intro-R-programming” folder as the `Project working directory`
6. Click `Open in new session`
7. Click `Create Project`

RStudio setup - optional

1. Choose **Tools** < **Global options**
2. Under **General**
 - DON'T Restore .RData into workspace at startup
 - NEVER Save workspace to .Rdata on exit:
3. Further personalize RStudio

Basic elements & data types

“To understand computations in R, two slogans are helpful: Everything that exists is an object. Everything that happens is a function call.”

— John Chambers

Basic elements & data types

- What are objects?
- Atomic vectors
- Vector structures
- Subsetting
- Replacement

What are objects?

- Data-structures that can be used in computations
- Collections of data of all kinds that are dynamically created and manipulated
- Can be very small, or very big. → *Everything in R is an object*
- Elementary data structures can be combined in more complex data structures
- Creating new types of *complex* objects is part of programming in R (S3, S4)

Atomic Vectors - Basic Building Blocks

Basic object types	
logical	TRUE, FALSE, NA
integer	1L, 142, -5, ..., NA
double	1.0, 1.25784, pi, ..., NA NaN, -Inf, Inf
character	"1", "Some other string", ..., NA

multiple values in one object → `length()` starting from 0

Atomic Vectors - Basic Building Blocks

Elements of the same type can be combined into an atomic vector using `c`.

```
c(3.3, 2.44, 9, 634)
```

```
> [1] 3.30 2.44 9.00 634.00
```

All elements are of the same type!

Atomic Vectors - Basic Building Blocks

An important object type with special behavior is **NULL**. It is an empty object that can be interpreted as *nothing*. It's length is 0.

```
length(NULL)
```

```
> [1] 0
```

NULL is mostly used as a default argument in functions, in order to create some default behavior.

Useful Functions

?seq Creates a vector with a sequence of numerical values.

```
seq(0, 10, by = 2)
```

```
> [1] 0 2 4 6 8 10
```

```
seq(0, 1, length.out = 11)
```

```
> [1] 0.0 0.1 0.2 0.3 0.4 0.5 0.6 0.7 0.8 0.9 1.0
```

Useful Functions

`seq_along` and `seq_len` are shortcuts.

```
seq_along(c("a", "b", "c", "d"))
```

```
> [1] 1 2 3 4
```

```
seq_len(10)
```

```
> [1] 1 2 3 4 5 6 7 8 9 10
```

Avoid : when programming!

Useful Functions

?**rep** Creates a new vector by repeating the elements of a vector.

```
rep(1:3, each = 2)
```

```
> [1] 1 1 2 2 3 3
```

```
rep(1:3, times = 2)
```

```
> [1] 1 2 3 1 2 3
```

Useful Functions

?rep Creates a new vector by repeating the elements of a vector.

```
rep(c("a", "b", "c"), times = 2)
```

```
> [1] "a" "b" "c" "a" "b" "c"
```

```
rep(c("this", "may", "be", "useful", "!"), 1:5)
```

```
> [1] "this"    "may"     "may"     "be"      "be"      "be"      "be"
> [9] "useful"  "useful"  "!"       "!"       "!"       "!"       "!"
```

Useful Functions

?paste Creates a character vector by pasting multiple vectors together.

```
paste("one", "big", "string", sep = " ")
```

```
> [1] "one big string"
```

```
paste0("word_", seq(1, 4))
```

```
> [1] "word_1" "word_2" "word_3" "word_4"
```

```
paste(c("ONE", "TWO"), seq(1, 3),  
      sep = " || ", collapse = "_-_" )
```

```
> [1] "ONE || 1_-_TWO || 2_-_ONE || 3"
```

Useful Functions

?unique Creates a vector with the unique values of a vector.

```
unique(c("b", "a", "a", "b"))
```

```
> [1] "b" "a"
```

Useful Functions

?sort Creates a sorted version a Vector.

```
sort(c("b", "a", NA, "a", "b"))
```

```
> [1] "a" "a" "b" "b"
```

```
sort(c("b", "a", NA, "a", "b"), na.last = TRUE)
```

```
> [1] "a" "a" "b" "b" NA
```

```
sort(c(4, 2, 6, 1, 3, 5), decreasing = TRUE)
```

```
> [1] 6 5 4 3 2 1
```

Coercion/Conversion

Automatic conversion:

NULL → logical → integer → double → character

```
1 + TRUE
```

```
> [1] 2
```

Explicit conversion:

```
as."type"() as.vector(, mode = "type")
```

```
as.logical(0:5)
```

```
> [1] FALSE TRUE TRUE TRUE TRUE TRUE
```

atomic vectors - check type

Check type using: `is."type"()`

```
is.null(NULL)
```

```
> [1] TRUE
```

Check type using: `typeof()`

```
typeof(TRUE + FALSE)
```

```
> [1] "integer"
```

Assignment

In order to compute with objects efficiently, names can be assigned to the objects using the assignment operator `<-` (or `=`)

```
my_object <- TRUE  
my_object  
  
> [1] TRUE
```

- The objects (with references) that are available to a user can be seen in the global environment using `ls()`.
- R overrides previous assignments without a message. Removed objects (`rm(objectName)`) cannot be restored.

→ *May the source code be with you!*

Attributes

Attributes can be attached to objects. An attribute:

- has a name
- is itself also an object
- attributes are easily lost in computations. (One of the reasons to use OOP with classes and methods.)

```
my_object <- structure(5,  
                        my_attribute = "string",  
                        other_attribute = FALSE)  
attributes(my_object)  
  
> $my_attribute  
> [1] "string"  
>  
> $other_attribute  
> [1] FALSE
```

Attributes

There are several attributes with a specific use: `"names"`, `"dim"`, `"class"`, `"levels"`

- `"names"` is a character vector that contains the names of elements of the vector/object. Names can be printed and set using `names(object) <-` .
- `"dim"` is an integer vector that specifies how we should interpret the vector (i.e., as a matrix, as an array). The dimensions of a vector can be printed and set using `dim(object) <-` .
→ a **matrix** or **array** is a vector with a `"dim"` attribute.

- `"class"` is a character vector that contains class names. Classes can be printed and set using `class(object) <- .`
See Object Oriented Programming (S3).
- `"levels"` is a character vector that contains the names levels of a factor. Levels can be printed and set using `levels(factor) <- .`

A factor in R is actually an integer vector with

- a **"class"** attribute set to **"factor"**
- a **"levels"** attribute set to the level-labels that correspond to the integer values from 1 to the highest integer value in the integer vector.

More Basic Object Types

More basic object types	
complex	<code>1 + 2.31i, ... NA</code>
raw	<code>as.raw(2), charToRaw("a")</code>
expression	<code>expression(1+1, sum(a, b))</code>
language	a function call, <code>quote(1 + y)</code>
closure	<code>function(x) x - 1, mean</code>
builtin	<code>sum, c</code>
special	<code>for, return</code>
environment	an environment
symbol	<code>quote(x)</code>
...	...

Vector Structures

More basic object types	
list	<code>list()</code> , <code>as.list()</code> , ...
matrix	a vector with "dim" argument: two dimensions <code>matrix()</code> , <code>as.matrix()</code> matrix algebra
array	a vector with with "dim" argument
data.frame	a list with vectors of equal length <code>data.frame()</code> , <code>as.data.frame()</code>

List

A list is a “vector” that can contain any type of elements

- the types of elements can differ \leftrightarrow atomic vectors
- possible elements including lists \rightarrow recursive
- can have attributes

```
my_list <- list("this",  
               a = list(a = c(1:2)))
```

```
my_list
```

```
> [[1]]  
> [1] "this"  
>  
> $a  
> $a$a  
> [1] 1 2
```

Matrix & Array

A matrix or an array is a vector with a "**dim**"-attribute

- mostly useful for numeric vectors (integer and double)
- matrix algebra! `t(matrix)`, `%*%`, `aperm(array)`, ...
- matrix has two dimensions, array has n dimensions You can create an matrix array using:
 - `cbind(vector1, vector2)`
 - `rbind(vector1, vector2)`
 - `matrix(vector, ncol = 4, nrow = 2)`
 - `array(vector, dim = c())`

A data.frame is a list of (named) vectors of equal length.

- has dimensions (but not a "**dim**"-attribute)
- the columns are the vectors
- the vectors can be lists (using **I()**).
- a data.frame has row names (but ignore these)

Subsetting - Atomic vectors

A subset of elements from a vector can be accessed using `object[selection]`, where `selection` is:

1. a **logical** vector with the same length of the original vector (TRUE: select; FALSE: don't select)
2. an **integer** vector indicating the indexes of the elements to select (or exclude)
3. a **character** vector with the names of the elements to select

Using a **logical** vector:

- the logical vector should have the same length as the object. If shorter, the logical is repeated; if longer, **NAs** are added if **TRUE**. → always use the same length!
- handy when you want to select based on a condition related to the object values

Subsetting - Atomic vectors

Using a **logical** vector:

```
my_object <- c(a = 1, b = 5, c = 3, d = 8)
my_object[my_object > 4]

> b d
> 5 8
```

Subsetting - Atomic vectors

Using an **integer** vector:

- the integer vector can have any length (repeated indices are repeatedly selected)
- positive values mean *select*, negative values mean *drop*
- positive and negative values cannot be combined
- for integers higher than the number of elements in the vector, **NAs** are added
- using **which()** a logical vector is transformed in an integer vector with the indices of the elements that were **TRUE**
- double elements are truncated towards zero (using **as.integer()**)

Subsetting - Atomic vectors

Using an **integer** vector:

```
my_object <- c(a = 1, b = 5, c = 3, d = 8)
my_object[c(1, 2, 1, 2, 1, 2, 1, 2, 1, 2, 1, 2)]
```

```
> a b a b a b a b a b a b
> 1 5 1 5 1 5 1 5 1 5 1 5
```

Subsetting - Atomic vectors

Using a **character** vector:

- the strings that match with the names of the elements in the vector are returned
- the character vector can have any length (repeated names are repeatedly selected)
- only selection is possible (dropping is not)
- strings that are not matched with names return **NA**

Subsetting - Atomic vectors

Using a **character** vector:

```
my_object <- c(a = 1, b = 5, c = 3, d = 8)
my_object[c("a", "c")]

> a c
> 1 3
```


Subsetting - Atomic vectors

A **single** element from a vector can be accessed using `object[[selection]]`, where **selection** is:

- an **integer** value indicating the index of the element to select
- a **character** vector with the name of the element to select

```
my_object <- c(a = 1, b = 5, c = 3, c2 = 8)
my_object[[2]]
```

```
> [1] 5
```

Subsetting - Matrix & Arrays

Because arrays and matrices are atomic vectors (with a `"dim"` argument), the rules for atomic vectors apply.

Subsetting - Matrix & Arrays

In addition, selection is possible per dimension:

- separated by a comma [,]
- selection via character (match row or column names), integer (row and column number) or logical vectors
- the first vector selects the rows, the second the columns (and so on)
- **dimensions are dropped**, unless `drop = FALSE`

```
my_matrix <- matrix(c(11, 21, 12, 22), ncol = 2,  
                    dimnames = list(paste0("row", 1:2),  
                                     paste0("col", 1:2)))  
my_matrix[,2]
```

```
> row1 row2  
>  12   22
```

Subsetting - Matrix & Arrays

Finally, the selection element can also be a matrix (with one column per dimension). Each row in the matrix selects one value.

```
my_matrix <- matrix(c(11, 12, 21, 22), ncol = 2,
                    dimnames = list(paste0("row", 1:2),
                                     paste0("col", 1:2)))
selection_matrix <- rbind(c(1, 1), c(1, 2), c(2, 1))
my_matrix[selection_matrix]

> [1] 11 21 12
```

Subsetting - Lists

For lists, the rules are similar as for atomic vectors.

- `list[selection]` gives a list (i.e., a subset of the original list)
- `list[[selection]]` gives the element (which can be a list)
- `list[["element_name"]]` is the same as `list$element_name`

```
my_list <- list(a = 1, b = 5, c = 3, d = 8)
is.list(my_list["a"])
```

```
> [1] TRUE
```

```
is.list(my_list[["a"]])
```

```
> [1] FALSE
```

Subsetting - data.frames

Because data.frames are lists, the rules for lists apply.

```
my_dat <- data.frame(col1 = c(11, 21),  
                     col2 = c(12, 22))
```

```
my_dat[1]
```

```
>   col1  
> 1   11  
> 2   21
```

Subsetting - data.frames

In addition, the selection rules for matrices can be used:

- selection per row and column (note the **drop** argument)
- selection via a matrix with two columns

```
my_dat <- data.frame(col1 = c(11, 21),  
                     col2 = c(12, 22))  
my_dat[, "col1", drop = FALSE]
```

```
>   col1  
> 1    11  
> 2    21
```

Subsetting - data.frame & matrixe

Programming advice

Code defensively: always use `, drop = FALSE`

Element Replacement

A subset of elements from a vector or vector structure can be replaced using `object[selection] <- new_values`:

- the modifications are done in place
- the structure and class of the object stay unchanged
- the length of the new values should correspond with the length of the selection (the number of elements to replace should be a multiple of the number of new values)
- *only for lists*: the replacement can be **NULL** (which removes the element from the list)

Element Replacement

```
my_dat <- data.frame(col1 = c(11, 21),  
                     col2 = c(12, 22))
```

```
my_dat[1, 2] <- 33
```

```
my_dat
```

```
>   col1 col2  
> 1   11   33  
> 2   21   22
```

Flow & conditional programming

Flow & conditional programming

R has specific tools (functions) that help organize the flow of computations.

You can make computations conditional on other objects (“conditional computation”)

The most commonly used tools are:

- `if (+ else)`
- `ifelse`

Conditional Computation - if

if statements have the basic form

```
if(test){  
    some_computations  
}
```

- **test** should be either **TRUE** or **FALSE** (or code that results in one of both).
- If **test == TRUE**, than **some_computations** is executed, if **test == FALSE**, than not.
- **Important:** **test** should have length 1. If not, only the first element is considered.

Conditional Computation - if

`else` can be added, but it is optional

```
if(test){  
    some_computations  
} else if (test_2){  
    other_computations  
} else {  
    more_computations  
}
```

Typical test functions

Vectorized, or elementwise	
<code>==</code>	equal to
<code>!=</code>	NOT equal to
<code>>, <</code>	is greater, less than
<code>>=, <=</code>	is greater, less than or equal to
<code>&</code>	AND operator
<code> </code>	OR operator
<code>xor</code>	exclusive OR

Typical test functions

Not Vectorized	
<code>identical()</code> <code>any()</code> <code>all()</code> <code>&&</code> <code> </code> <code>is.character()</code> , <code>is.data.frame()</code> , ...	identical to at least one TRUE all TRUE AND operator OR operator

Typical test functions

Compare:

```
c(TRUE, TRUE) & c(FALSE, TRUE)
```

```
> [1] FALSE TRUE
```

```
c(TRUE, TRUE) && c(FALSE, FALSE)
```

```
> [1] FALSE
```

Conditional Computation - if

The *test* should have length 1!

```
# only the first element is evaluated
age <- c(8, 17, 39, 55)
if (age >= 18) {
  "can vote"
} else {
  "too young"
}
```

```
> Warning in if (age >= 18) {: the condition has length
> 1 and only the first element will be used
```

```
> [1] "too young"
```

Conditional Computation - if

Typical uses

```
if(any(is.na(x))){  
  stop("computation impossible due to NA values")  
}  
  
if(!is.integer(vector)){  
  warning("'vector' is automatically converted to interger.  
    This may affect the results")  
  vector <- as.integer(vector)  
}  
  
if(is.null(argument)){  
  # default computations  
} else if (argument == specific_value) {  
  # other computations  
}
```

Programming advice

- *if* is almost always used inside of functions or loops
- If possible, avoid using *else*
- Use meaningful initialization, early return(), stop(), etc. instead

Conditional Computation - if

Solution using *if* and *else*

```
age <- 17
if (age >= 18) {
  vote <- "can vote"
} else {
  vote <- "too young"
}
vote

> [1] "too young"
```

Conditional Computation - if

Solution using meaningful initialization

```
age <- 17
vote <- "too young"
if (age >= 18) {
  vote <- "can vote"
}
vote

> [1] "too young"
```

Conditional Computation - ifelse

A vectorized version is `ifelse()`.

```
# all elements are evaluated
age <- c(8, 17, 39, 55)
ifelse(age >= 18,
       yes = "can vote",
       no  = "too young")

> [1] "too young" "too young" "can vote"  "can vote"
```

Conditional Computation - ifelse

Go-to tool for conditional recoding

```
age_estimated <- c(10, 20, 35, 60)
age_self_rep <- c(NA, 17, 39, NA)

# Use available information, prioritize self report
ifelse(!is.na(age_self_rep),
      yes = age_self_rep,
      no = age_estimated)

> [1] 10 17 39 60
```


Exercises



Loops & Iteration

Loops & iteration

R has specific tools (functions) that help organize the flow of computations.

You can repeat a similar computation multiple times typically with changing options (“iteration”). The most commonly used tools are:

- loops (**repeat**, **while**, **for**)
- functionals (**apply** - family)

Loops & Iteration - for

for statements have the basic form

```
for (element in vector) {  
    computation  
}
```

For each element in the vector, the computation is executed. Often, the computation depends on the element in that iteration.

Loops & Iteration - for

```
for (index in 1:3){  
  cat(" computation -")  
}
```

```
> computation - computation - computation -
```

```
for (name in c("Alice", "Bob", "Casey")){  
  if(name == "Bob") cat(" This was Bob -")  
  else cat(" Not Bob -")  
}
```

```
> Not Bob - This was Bob - Not Bob -
```

Loops & Iteration - for

Nested loops (over the rows and columns of a matrix)

```
matrix <- matrix(NA, nrow = 2, ncol = 3)
for (rowNr in 1:2){
  for (colNr in 1:3){
    matrix[rowNr, colNr] <- rowNr * 10 + colNr
  }
}
matrix
```

```
>      [,1] [,2] [,3]
> [1,]   11   12   13
> [2,]   21   22   23
```

Loops & Iteration - while

`while` statements have the basic form

```
while (condition){  
    computation  
}
```

As long as the condition is TRUE, the computation is executed. Often, the computation depends on something that is related to the condition.

Loops & Iteration - while

Sample five random values from a normal distribution, the distance between the minimum and maximum should be at least 4.

```
max_dif <- 0
while (max_dif <= 4){
  cat("|")
  values <- rnorm(5)
  max_dif <- max(values) - min(values)
}
max_dif
round(values, 3)
```


Loops & Iteration - repeat

`repeat` statements have the basic form

```
repeat {  
  computation  
}
```

Without a `break` the computation is repeated infinite times

Loops & Iteration - next break

- **next** starts next iteration
- **break** ends iteration (of the innermost loop)

```
index <- 0
repeat {
  index <- index + 1
  if (index %in% c(3, 5)) next
  if (index > 6) break
  print(index)
}
```

```
> [1] 1
> [1] 2
> [1] 4
> [1] 6
```

Iteration - Good practice

Programming advice

Use `seq()`, `seq_len()`, or `seq_along()`.

```
x <- numeric()
for (index in 1:length(x)){
  print(index)
}

> [1] 1
> [1] 0

for (index in seq_along(x)){
  print(index)
}
```

Loops & Iteration - Good practice

Programming advice

Don't grow, replace.

```
x <- letters
result1 <- numeric()           # grow
result2 <- numeric(length(x)) # replace
for (index in seq_along(x)){
  result1 <- c(result1, paste(index, x[index])) # grow
  result2[index] <- paste(index, x[index])      # replace
}
```

Loops & Iteration - Functionals

A functional is a function that takes another function as an argument.

Focus on the **apply**-family. These functions *apply* a function repeatedly.

Can be seen as an abstraction of a for loop, with the following advantages

- requires less code to write
- does not store intermediate results
- no need to replace / grow

Functionals

The most commonly used functionals are:

- `lapply` vector / list \rightarrow list
- `sapply` vector / list \rightarrow vector (matrix)
- `apply` matrix / array / data.frame \rightarrow vector (matrix)
- `tapply`, `by`, `aggregate`
- `mapply`, `Map`
- `rapply`, `eapply`, `vapply`

All of which have an argument that should be a function.

lapply

data.frames are lists with the columns as elements:

```
lapply(iris, FUN = class)
```

```
> $Sepal.Length
```

```
> [1] "numeric"
```

```
>
```

```
> $Sepal.Width
```

```
> [1] "numeric"
```

```
>
```

```
> $Petal.Length
```

```
> [1] "numeric"
```

```
>
```

```
> $Petal.Width
```

```
> [1] "numeric"
```

```
>
```

```
> $Species
```

```
> [1] "factor"
```

lapply

- any type of element can be used
- other arguments can be passed through

```
lapply(airquality, FUN = mean, na.rm = TRUE)
```

```
> $Ozone
```

```
> [1] 42.12931
```

```
>
```

```
> $Solar.R
```

```
> [1] 185.9315
```

```
>
```

```
> $Wind
```

```
> [1] 9.957516
```

```
>
```

```
> $Temp
```

```
> [1] 77.88235
```

```
>
```

```
> $Month
```

```
> [1] 6 2024/6/
```


apply

- for objects with dimension (matrix, array, data.frame)
- apply over (a) chosen dimension(s)

```
my_matrix <- matrix(1:6, nrow = 2)
apply(my_matrix, 1, max)      # apply per row

> [1] 5 6

apply(my_matrix, 2, max)      # apply per column

> [1] 2 4 6
```

apply

```
my_array <- array(1, dim = c(2, 3, 4))
apply(my_array, c(1, 2), sum)      # apply per row and column

>      [,1] [,2] [,3]
> [1,]    4    4    4
> [2,]    4    4    4

apply(my_array, 3, sum)            # apply per "third dimension"

> [1] 6 6 6 6
```

Exercises



Functions I

“To understand computations in R, two slogans are helpful: Everything that exists is an object. Everything that happens is a function call.”

— John Chambers

Function Calls

Computing in R happens through function calls. A function is applied to one or more objects, and returns an object after the computation.



Figure 1: A function call.

The typical use is:

```
function(object1, argument = object2)
```

- Computations that seem not to be done using function calls are actually also function calls. Try ``<`-(a, 5)` or ``>`(5, 2)`
- most functions that seem not to return an object, return it invisibly. Check `(a <- 5)`.

Building Blocks

Functions are the building blocks of R code. Writing functions allows you to organize and optimize the computations that you want to do.

Functions should:

- have a clear purpose
- be well documented
- be portable

Central stepping stone for R users:

Move from solely using functions written by others to writing your own functions.

Function definition

- Name
- Arguments/Formals (input)
- Body (what happens inside, R-code with the computations)
- Output

Function definition

```
countNA <- function(x) {  
  out <- sum(is.na(x))  
  out  
}
```

Name
Arguments/Formals
Body
Output

Function Names

Every function needs a (meaningful) name!

- Usually **a verb** (what does the function do?)
- Avoid existing names
- Better longer than unclear
- CamelCase vs snake_case

Function Names

Good

- `computeAIC()`
- `removeNAs()`
- `drop_NA_rows()`
- `factor_to_dummies()`

Bad

- `myFun()`
- `foo()`
- `statistics()`
- `data_preparation()`

Most functions take one or multiple inputs.
These are usually:

- One or two data arguments
- Additional Options

Examples for zero arguments

- `getwd()`
- `Sys.time()`

Examples for one argument

- `dim()`
- `names()`

Examples for multiple arguments

- `mean()`
- `median()`
- `lm()`

Programming advice

Less arguments = better!

Functions usually return a single object, namely the last evaluated object.

```
get_log_xtox <- function(x) {  
  x_x <- x^x  
  out <- log(x_x)  
  out  
}  
get_log_xtox(2)
```

Often arguments have to be objects of a specific type.

```
sum(c("a", "b", "c")) # gives an error
```

The documentation typically gives (or should give) information about what objects the arguments should be. Check `?sum`

Exercises



That's it for today!

That's it for today!

Questions? Remarks?